

Query Optimization Strategy

To ensure **scalability** and **high performance** for subscription data access, several targeted optimizations were implemented:

1. Raw SQL for Performance-Critical Paths

- I used **SQLAlchemy Core with raw SQL** instead of ORM `.query.filter()` for subscription listing and history retrieval.
- This avoids SQLAlchemy's overhead and gives full control over the query structure.

```
SELECT id, plan_id, start_date, end_date
FROM user_subscription
WHERE user_id = :user_id
ORDER BY start_date DESC
```

2. Indexing Strategy

Indexes were carefully added to **match query filters and sorting** conditions:

Index Name	Columns	Purpose
idx_user_active	(user_id, is_active)	Speeds up lookup for active subscriptions
idx_user_start_date	(user_id, start_date)	Optimizes order-by for listing history
idx_user_plan	(user_id, plan_id)	Supports filtering or joining by plan
Recommended	end_date	Needed for fast access to historical data

3. Pagination with offset-limit

- Endpoints that return lists (e.g., active subscriptions, historical subscriptions) use LIMIT and OFFSET to reduce result set size and improve latency:

```
LIMIT :page_size OFFSET :offset
```

- Separate total count queries: Count queries are isolated from data-fetching queries to improve clarity and avoid unnecessary joins:

```
SELECT COUNT(*) FROM user_subscription WHERE user_id = :user_id AND is_active = TRUE
```

4. Bulk Insert/Seed Optimization

- Data seeding is done using **batched inserts** to avoid per-row overhead.
- Created foreign key references for `user` and `plan` before seeding subscriptions to maintain integrity.

5. Query Performance Testing

- Included benchmark tests using `pytest` to simulate querying against a database with **50,000+ rows**.
- Visual comparisons (see `images/`) demonstrate the performance gains when indexes are in place.

Why `end_date` IS NOT NULL?

- Used to separate **historical subscriptions** from active ones.
- This condition is selective; indexing `end_date` improves speed significantly when querying history.

Pagination Example in Active Subscriptions

- Total count and paginated results are queried **separately** for flexibility and performance.
- This approach supports scalable frontend pagination, load more, or infinite scroll.

Repository layer

```
get_active_subscriptions_paginated(user_id, page, page_size)
count_active_subscriptions(user_id)
```

Service layer

```
{
    "subscriptions": [...],
    "total": 128,
    "page": 2,
    "page_size": 10
}
```

ORM Overhead in Subscription Handling

Overhead in ORM Queries

- **Session Management:** ORM maintains a session and tracks object states (e.g., dirty-checking), which adds memory and CPU usage.
- **Object Construction:** Each database row is converted into a Python object, even if I only need raw data.
- **Lazy Relationships:** ORM loads relationships lazily by default, which can result in **N+1 query problems** unless explicitly optimized.
- **Join Complexity:** ORM-generated SQL for relationships can be verbose and suboptimal for large datasets.

Example (inefficient with ORM):

```
subscriptions =
UserSubscription.query.filter_by(user_id=1).order_by(UserSubscription.start_date
.desc()).all()
```

- Loads full ORM models into memory.
- Triggers unnecessary relationship loading unless optimized.
- Harder to control SELECT columns and query plan.

Optimization Decision

In this project:

- We used **SQLAlchemy Core with raw SQL** for:
 - `list_all_subscriptions`
 - `get_subscription_history`
- This bypasses the ORM session and constructs only what we need—**rows, not objects**.

Result:

- Lower latency for read-heavy endpoints.
- Faster data access and reduced memory usage, especially when handling **50K+ subscriptions**.